
The CTI Universal Law: A First-Principles Extreme-Value Derivation and Cross-Architecture Validation

Anonymous authors

Paper under double-blind review

Abstract

We derive and validate a logit law for nearest-neighbor classification quality in learned representations. Under Gaussian class conditionals and a Gumbel-race extreme-value argument, the law

$$\text{logit}(q_{\text{norm}}) = \alpha \kappa_{\text{nearest}} - \beta \log(K - 1) + C_0$$

($q_{\text{norm}} = (\text{acc}_{1\text{NN}} - 1/K) / (1 - 1/K)$, $\kappa_{\text{nearest}} = \min_{j \neq k} \|\mu_j - \mu_k\| / (\sigma_W \sqrt{d})$) is a *conditional theorem*: the functional form follows from first principles; only constants (α, β, C_0) require estimation.

Leave-one-architecture-out (LOAO) over 12 NLP architectures yields $\hat{\alpha} = 1.477 \pm 0.034$ with $\text{CV} = 0.023$ ($R^2 = 0.955$). Cross-modal validation confirms the same form for ViT-Large ($R^2 = 0.964$), ResNet50, and 7 frozen speech models on Speech Commands ($r = 0.898$, $p = 0.006$, $K = 36$). A pre-registered causal confusion-matrix prediction ($r = 0.842$ – 0.776 , sign accuracy 92.9–100%, $n = 182$ pairs, all $p < 10^{-35}$) and a strictly blind OOD test ($r = 0.817$, $p = 0.013$) establish the law’s predictive power. Biological generalization: 30/32 mouse visual-cortex sessions satisfy the law ($\bar{r} = 0.736$). The Gumbel-race mechanism extends to generation: κ from the unembedding matrix W_U alone predicts log CE across 22 models ($r = -0.55$, $p = 0.008$; $r = -0.70$ excluding one outlier).

Honest scope: α varies by signal type (NLP decoders: 1.48; continuous-signal modalities: ViT 4.5, CNN 4.0, audio 4.7); universality is of the *functional form* only.

1 Introduction

Neural scaling laws established simple macroscopic regularities for loss as a function of compute and model size (Kaplan et al., 2020; Hoffmann et al., 2022), though unreliable for downstream tasks (Lourie et al., 2025). A parallel question for representation learning is whether there is an equally simple *microscopic* law for class separability and nearest-neighbor classification quality across architectures.

This paper presents the CTI Universal Law:

$$\text{logit}(q_{\text{norm}}) = \alpha \kappa_{\text{nearest}} - \beta \log(K - 1) + C_0. \quad (1)$$

The key conceptual point: Eq. 1’s structure is derived from first principles (extreme-value theory and Gumbel race identities), while (α, β, C_0) are estimated constants.

Contributions.

1. **Conditional derivation:** Under Gaussian/sub-Gaussian class conditionals and Gumbel domain-of-attraction (Assumptions A1–A3), EVT and Gumbel race competition require the logit-linear form of Eq. 1.
2. **12-architecture LOAO test** (pre-registered): LOAO $\text{CV}_\alpha = 0.023$ ($R^2 = 0.955$) spanning transformers, SSM hybrid, and pure linear RNN (RWKV-4). Extended to 19 architectures with $\text{CV}_\alpha = 1.75\%$.

- 37 3. **Cross-modal and causal validation:** ViT-Large $R^2 = 0.964$; ResNet50 CNN $r =$
38 0.749 ; 7 speech models $r = 0.898$ ($K = 36$); pre-registered confusion-matrix causal
39 prediction $r = 0.842$ – 0.776 with zero refit; blind OOD $r = 0.817$; expanded holdout
40 (11 models $\times$ 8 datasets, $n = 77$) all 6 criteria PASS ($r = 0.879$, MAE= 0.077).
41 4. **Biological validation:** 30/32 Allen Neuropixels mouse V1 sessions PASS; multi-area
42 batch confirms $r > 0.70$ across 5 cortical areas.
43 5. **Generation law:** κ from W_U (zero forward passes) predicts log CE across 22 models
44 ($r = -0.55$; $r = -0.70$ excluding one outlier; fixed- V : $r = -0.84$, architecture-
45 independent). Vocabulary size drops out ($K_{\text{eff}} \approx 2$ – 3).

46 The law enables zero-cost diagnostics: κ_{nearest} from frozen embeddings ranks architectures
47 by retrieval quality ($\rho = 0.833$, $p = 0.005$; pre-registered) and selects optimal representation
48 layers without downstream labels. Four probe measurements reduce MAE by 86%.

49 2 Related Work

50 **Neural scaling laws.** Power-law regularities for language-model loss are well-established
51 (Kaplan et al., 2020; Hoffmann et al., 2022). Our aim is analogous but targets geometry-to-
52 accuracy transfer within individual models.

53 **Representation quality and neural collapse.** kNN-based quality diagnostics are widely
54 used (Wu et al., 2018; Caron et al., 2021). Neural collapse geometry (Papayan et al., 2020;
55 Han et al., 2022; Zhu et al., 2021; Rangamani et al., 2023; Wu and Papayan, 2024) and
56 spectral order parameters (Garrido et al., 2023; Liu et al., 2025; Munn et al., 2024) predict
57 downstream performance. None derive a universal slope connecting κ_{nearest} to $\text{logit}(q)$
58 across architectures.

59 **EVT and Gumbel races.** Classical EVT describes minima/maxima asymptotics (Fisher
60 and Tippett, 1928; Gnedenko, 1943; Gumbel, 1958); Zhang et al. (2025) independently
61 confirm Gumbel priors for nearest-neighbor matching. Our derivation applies this to
62 distance minima among competing classes. For generation, Wu and Papayan (2024) find NC
63 properties correlate with generalization in LLMs; Kulkarni et al. (2026) show effective rank
64 alone does not predict loss; our κ outperforms by capturing competition-relevant geometry.

65 3 Theory: Derivation of the CTI Universal Law

66 3.1 Definitions

67 Let K be the number of classes, d the embedding dimension, and $\text{acc}_{1\text{NN}}$ the 1-NN accuracy.
68 The chance-normalized accuracy is $q_{\text{norm}} = (\text{acc}_{1\text{NN}} - 1/K) / (1 - 1/K)$. For class centroids
69 $\{\mu_k\}_{k=1}^K$, pooled within-class per-dimension std σ_W , and embedding dimension d :

$$\kappa_{\text{nearest}}(k) = \min_{j \neq k} \frac{\|\mu_j - \mu_k\|}{\sigma_W \sqrt{d}}. \quad (2)$$

70 The denominator $\sigma_W \sqrt{d}$ equals the expected within-class L2 radius, making κ_{nearest} a di-
71 mensionless signal-to-noise ratio.

72 3.2 Gumbel race mechanism

73 Consider a query x from class k . Let $M_+ = \min_{i: y_i = k, i \neq x} \|x - x_i\|$ and $M_j = \min_{i: y_i = j} \|x -$
74 $x_i\|$ for $j \neq k$. 1-NN is correct iff $M_+ < \min_{j \neq k} M_j$.

75 **Proposition 1** (Gumbel race identity). *If $G_0, G_1, \dots, G_{K-1}$ are independent Gumbels with*
76 *common scale s and locations $\mu_0, \dots, \mu_{K-1}$, then $\Pr(G_0 < \min_{j \geq 1} G_j) = [1 + \sum_{j \geq 1} \exp(-(\mu_j -$
77 $\mu_0)/s)]^{-1}$.*

78 Applying this with $G_0 \equiv M_+$ gives

$$q_{\text{norm}} \approx \left[1 + \sum_{j \neq k} \exp\left(-\frac{\Delta_{k,j}}{s}\right) \right]^{-1}, \quad (3)$$

79 where $\Delta_{k,j} := \mu_{k,j}^{(-)} - \mu_k^{(+)}$ is the gap between cross-class and same-class nearest-neighbor
80 distances.

81 3.3 From Eq. 3 to Eq. 1

82 **Nearest-class dominance.** For equidistant competitors: $\sum_{j \neq k} \exp(-\Delta_{k,j}/s) \approx (K -$
83 $1) \exp(-\Delta_k/s)$, giving $\beta = 1$ asymptotically. Real datasets show hierarchical class structure
84 with $\beta \approx 0.5$ (sparse competition, $N_{\text{eff}} \propto \sqrt{K} - 1$; see Appendix E).

85 **Linear gap-to-SNR map.** Under isotropic geometry, $\Delta_k = a_0 + a_1 \kappa_{\text{nearest}}(k) + o(1)$, yield-
86 ing:

87 **Theorem 2** (CTI logit law: functional form). *Under (A1) sub-Gaussian class conditionals*
88 *with shared scale, (A2) Gumbel domain-of-attraction for within-class distances, (A3) nearest-class*
89 *dominance, as $n, d \rightarrow \infty$ with K fixed:*

$$q_{\text{norm}} = \sigma(\alpha \kappa_{\text{nearest}} - \beta \log(K - 1) + C_0) + \varepsilon_{d,n}, \quad \varepsilon_{d,n} \rightarrow 0.$$

90 *The functional structure is derived; only (α, β, C_0) are estimated. Deviations produce measurable*
91 *residuals (§7).*

92 **Retracted equicorrelation analysis.** The former universal-equicorrelation, near-simplex-
93 universality, and zero-parameter $\alpha(\rho)$ claims are retracted; Appendix A records the historical
94 measurement and the failed survival tests.

95 4 Experimental Design

96 4.1 Primary 12-architecture LOAO test

97 We evaluate 12 architectures on 4 text classification datasets: AGNews ($K = 4$), DBpedia
98 ($K = 14$), 20Newsgroups ($K = 20$), and GoEmotions ($K = 28$). For each model we extract
99 embeddings at 4 proportional-depth layers, yielding 192 total points. Architectures span
100 2021–2025: GPT-NeoX (Pythia-160m/410m/1b), GPT-Neo-125m, Qwen2.5-0.5B, Qwen3-
101 0.6B/1.7B, OLMo-1B, TinyLlama-1.1B, Mistral-7B, Falcon-H1-0.5B (hybrid), and RWKV-4-
102 169m (pure linear RNN). Pre-registered thresholds: $\text{CV}_\alpha < 0.25$, $\text{CV}_\beta < 0.25$.

103 5 Results

104 5.1 Global fit and LOAO constants

105 With per-dataset intercepts: $\hat{\alpha} = 1.477$, $\hat{\beta} = 0.326$, $R^2 = 0.955$.

106 $\text{CV}_\alpha = 0.019$ (single- C_0) and 0.023 (per-dataset) are both $10\times+$ below the pre-registered 0.25
107 threshold. RWKV-4 (pure linear RNN, no attention) satisfies the pre-registered boundary
108 test: $\hat{\alpha} = 2.887 \in [2.43, 3.29]$. See Figure 1.

109 5.2 Cross-modal validation

110 ViT-Large at $K = 10$ reaches $R^2 = 0.964$, matching the NLP fit in explained variance. The
111 ViT CIFAR-10 entries report A (not α): they fit $\text{logit}(q) = A \cdot \kappa \cdot \sqrt{d_{\text{eff}}} + C$, factoring out the

Table 1: 12-architecture LOAO estimates (single- C_0 fit, $R^2 = 0.627$). Per-dataset-intercept fit gives $\hat{\alpha} = 1.477$, $CV = 0.023$, $R^2 = 0.955$.

Architecture	Family	α_{LOAO}	β_{LOAO}	MAE
Pythia-160m	GPT-NeoX (2021)	2.800	0.758	0.710
Pythia-410m	GPT-NeoX (2021)	2.831	0.751	0.639
Pythia-1b	GPT-NeoX (2021)	2.847	0.750	0.664
GPT-Neo-125m	GPT-Neo (2021)	2.848	0.742	0.651
Qwen2.5-0.5B	Qwen (2024)	2.944	0.733	0.641
OLMo-1B	OLMo (2024)	2.926	0.733	0.668
TinyLlama-1.1B	LLaMA (2024)	2.808	0.750	0.641
Qwen3-0.6B	Qwen3 (2025)	2.943	0.745	0.619
Qwen3-1.7B	Qwen3 (2025)	2.943	0.734	0.605
Mistral-7B-v0.3	Mistral (2023)	2.840	0.738	0.786
Falcon-H1-0.5B	Hybrid (2025)	2.805	0.763	0.547
RWKV-4-169m	Linear RNN (2022)	2.887	0.751	0.692
Mean $\pm$ std	—	2.868 ± 0.055	0.746 ± 0.010	—
CV	—	0.019	0.013	—
Result	—	PASS	PASS	—




Figure 1: Left: Observed vs. predicted $\text{logit}(q_{\text{norm}})$ for 192 NLP points across 12 architectures ($R^2 = 0.955$). RWKV ($\star$) and Falcon-H1 ($\blacktriangle$) highlighted. Right: LOAO α per architecture; all within 6% of the mean.

112 effective dimensionality explicitly.¹ The functional form (logit-linear in κ) is universal; the
 113 constant is signal-type-specific.

114 **Audio modality.** Seven frozen speech models from 4 architectures (Wav2Vec2, HuBERT,
 115 WavLM, Whisper; 39M–316M parameters) on Speech Commands v0.02 ($K = 36$ spoken
 116 keywords, 4890 test samples; features at $\frac{2}{3}$ -depth layer, PCA to 256 dimensions) yield
 117 $r = 0.898$ ($p = 0.006$), $\alpha_{\text{audio}} = 4.669$ ($R^2 = 0.807$; supplementary materials). The slope
 118 $\alpha_{\text{audio}} \approx 4.7$ matches $\alpha_{\text{CNN}} \approx 4.4$ but not $\alpha_{\text{NLP}} \approx 1.5$, suggesting a two-regime structure:
 119 continuous-signal modalities (vision, audio) share $\alpha \approx 4$ –5, while discrete-token modalities
 120 (NLP) operate at $\alpha \approx 1.5$.

¹The CIFAR-10 ViT entries report A from $\text{logit}(q) = A \cdot \kappa \cdot \sqrt{d_{\text{eff}}} + C$, not α from $\text{logit}(q) = \alpha \cdot \kappa + C$. When ViT is fitted with the standard parameterization, the combined cross-layer $\alpha=4.49$, consistent with CNN (4.0) and audio (4.67).

Table 2: Cross-modal results. α is the fitted slope; Pearson r within each setting.

Model	Family	Dataset	K	α (A)	R^2/r
ViT-L-16-224	ViT	CIFAR-10	10	$A=0.630^\ddagger$	$R^2=0.964$
ViT-B-16-224	ViT	CIFAR-100	100	3.899	$r=0.773$
ResNet50	CNN	CIFAR-100	100	4.418	$r=0.749$
Audio (7 models)	4 arch families	Speech Cmds	36	4.669	$r=0.898$
NLP reference	Trans/SSM/RNN	4 datasets	4–28	1.477	$R^2=0.955$

Table 3: Causal evidence summary. Extended details in Appendix D.

Statistic	Value	Status
Arm A $r(\kappa_{j1}, \text{logit } q)$	0.899	Nearest drives logit
Arm C $r(\kappa_{jK}, \text{logit } q)$	0.000	Negative control PASS
Dose-response (6 unseen archs)	6/6 directional	Cross-modal PASS
Global surgery ratio / d_{eff}	18.74/24.52	PASS ($\pm 15\%$)
Confusion-matrix r ($\delta = 1$, no refit)	0.842	PASS (> 0.50)
Confusion-matrix r ($\delta = 2$, no refit)	0.808	PASS (> 0.50)
Confusion-matrix r ($\delta = 3$, no refit)	0.776	PASS (> 0.50)
Confusion-matrix sign acc. ($\delta = 2, 3$)	100%	PASS ($> 80\%$)

5.3 External blind OOD prediction and expanded holdout

Blind OOD. SmolLM2-1.7B (absent from training) with all parameters frozen: Pearson $r = 0.817$, $p = 0.013$ (**PASS**, threshold $r > 0.80$).

H8+ expanded holdout. 11 unseen models $\times$ 8 datasets ($n = 77$), all parameters frozen (pre-registered before data collection). All six criteria pass: $r = 0.879$ (H8a: ≥ 0.85), MAE = 0.077 (H8b: ≤ 0.10), 85.7% within 0.15 (H8c: $\geq 80\%$), per-family r : encoder = 0.948, decoder = 0.859 (H8d: ≥ 0.70), mean signed error = +0.041 (H8e), outlier rate = 3.9% (H8f: $\leq 10\%$). CTI’s contribution is *within-dataset architecture ranking* ($\bar{\rho} = 0.642$ on 4 high-variation datasets vs. $\rho = 0$ for dataset-mean baseline). See Figure 2 and Appendix G.

Cross-family transfer (LOMFO). Leave-one-model-family-out removes *all* models from a family, refits from remaining families, and predicts the excluded family. All four families achieve $r \geq 0.84$ (encoder: 0.943; decoder: 0.846; SSM: 0.984; hybrid: 0.959), confirming logit-linear structure transfers across families even with a misspecified slope. LODO (leave-one-dataset-out) mean $r = 0.125$ confirms that C_d requires per-dataset calibration.

5.4 Causal evidence

Three tiers of causal evidence establish κ_{nearest} as the causal driver: (i) frozen do-interventions and orthogonal factorial (nearest-arm $r = 0.899$, negative control $r = 0.000$); (ii) cross-modal dose-response (6/6 directional PASS, MAE = 0.0104); (iii) pre-registered causal confusion-matrix prediction with *fixed* $A = 1.054$, $\tau^* = 0.20$ and zero refit: $r = 0.842$ –0.776, sign accuracy 92.9–100%, $n = 182$ class pairs at each of three shift levels (all $p < 10^{-35}$). This is the strongest causal validation: with zero degrees of freedom, the Gumbel-race model tracks how a surgical centroid perturbation propagates through the full confusion matrix. See Appendix D for the full competition field analysis.

5.5 Biological validation

Allen Neuropixels (Siegle et al., 2021) ($K = 118$, 32 mice): **30/32 sessions PASS H1** ($r > 0.50$); mean = 0.736, CV = 15.4%; two non-passing sessions explained by noise-floor and ceiling regimes. Pre-registered multi-area batch (30 mice): VISp 30/30 (100%), VISl 22/22 (100%),




Figure 2: H8+ expanded holdout: predicted vs. actual q_{norm} for 11 unseen architectures across 8 datasets ($n = 77$). All parameters frozen. Encoder (blue) and decoder (orange) shown separately.

148 VISal 19/20 (95%), VISam 24/25 (96%), VISrl 20/23 (87%). All 4 secondary areas pass
 149 H_area3 ($\geq 87\%$). Convergent external validation: [St-Yves et al. \(2025\)](#) independently
 150 derived the same centroid-separation geometry in human visual cortex (fMRI, $K = 50$).

151 **Retracted competition-geometry analysis.** The earlier equicorrelation and near-simplex
 152 interpretations of the Allen data used the same incorrect whitening transform and are
 153 retracted. The biological validation here concerns the association between κ_{nearest} and
 154 decoding performance; it does not support a universal equicorrelation parameter.

155 5.6 Extension to next-token prediction (generation law)

156 Next-token prediction is V -way classification at each position. The CGF predicts that κ
 157 from W_U should predict cross-entropy, since the Gumbel-race mechanism depends only on
 158 output-space competition. We extract κ_{top1K} (nearest-neighbor distance among top-1000
 159 frequent token embeddings in W_U) from 22 models spanning Transformers ($n = 12$), SSMs
 160 (Mamba, $n = 5$), Hybrids ($n = 4$), and LiquidAI ($n = 1$).

161 **Key results:** $r(\kappa_{\text{top1K}}, \log \text{CE}) = -0.697$ ($p < 0.001$, $n = 21$ excluding LFM2.5 outlier).
 162 Partial r controlling for $\log N_{\text{params}} = -0.546$ ($p = 0.013$): not a model-size proxy. Fixed- V
 163 group (Pythia + Mamba, $V = 50280$, $n = 10$): $r = -0.837$ ($p = 0.003$), architecture-
 164 independent (F -test $p = 0.147$). Fitting $\log \text{CE} = -\alpha_{\text{gen}}\kappa + \beta_{\text{gen}}\log(V - 1) + C$ yields
 165 $\beta_{\text{gen}} \approx 0$: vocabulary size drops out because softmax competition reduces to $K_{\text{eff}} \approx 2\text{--}3$. See
 166 Figure 4 and Appendix I.




Figure 3: Allen Neuropixels biological validation ($K = 118$, 32 mice). Left: Pearson r per session; 30/32 PASS H1 ($r > 0.5$). Right: Two non-passing sessions explained by regime effects.




Figure 4: Generation law: κ from W_U (zero forward passes) predicts $\log(\text{CE})$. Left: All models ($n=22$, $r=-0.55$; regression line fit on $n=21$ excluding LFM2.5: $r=-0.697$). Right: Fixed-V (Pythia+Mamba, $n=10$; $r=-0.837$, architecture-independent).

5.7 Scope of universality and architecture ranking

The law exhibits a three-level universality hierarchy: **(1) Form**: logit-linear in κ_{nearest} , confirmed across NLP, ViT, CNN, audio, and biological systems. **(2) Constant within signal type**: α is stable within families (NLP decoders: $\text{CV} = 0.023$; 19 architectures: $\text{CV} = 1.75\%$); across families, α varies by signal type: discrete-token NLP decoders give $\alpha \approx 1.48$, while all continuous-signal modalities converge at $\alpha \approx 3.3\text{--}5.0$ (ViT: 4.5, CNN: 4.0, audio: 4.7). **(3) Intercept is task-specific**: C_d requires per-dataset calibration (LODO mean $r = 0.125$).

Practical architecture ranking (H3, pre-registered). A pre-registered test evaluates whether κ_{nearest} alone can rank 9 architectures by MAP@10 retrieval quality on Banking77 ($K = 77$): Spearman $\rho = 0.833$, $p = 0.005$ (PASS). κ_{nearest} is computed from geometric structure alone—no downstream labels, no probing, no retrieval run.

6 Discussion

What is universal. Seven strongly supported claims: (1) form universality (EVT-derived), (2) NLP constant stability ($\text{CV}_\alpha = 0.019\text{--}1.75\%$ across 12–19 architectures), (3) architecture boundary crossing (RWKV pure RNN passes), (4) cross-family form universality (ViT and

CNN yield comparable r at $K = 100$), (5) sparse competition exponent ($\beta = 0.5$ beats $\beta = 1.0$ by 10.1% LOAO; unconstrained $\hat{\beta} = 0.478$), (6) generation extension ($r = -0.55$ across 22 models, $r = -0.70$ excluding one outlier; architecture-independent within fixed V), (7) audio-visual alpha convergence: speech models from 4 architectures ($\alpha_{\text{audio}} = 4.67$, $r = 0.898$, $p = 0.006$) match vision CNNs ($\alpha_{\text{CNN}} = 4.4$), revealing a two-regime structure: continuous-signal modalities share $\alpha \approx 4\text{--}5$, while discrete-token NLP operates at $\alpha \approx 1.5$.

Honest scope. Universality is of the functional form; constants are family-specific. Encoders with CLS pooling fail ($r = 0.16$); with mean-pool, $\alpha_{\text{encoder}} \approx 7.1$. Encoder LOAO $\text{CV}_\alpha = 0.42$ (FAIL): encoder α is not family-invariant. C_d encodes dataset-specific difficulty requiring calibration. 4 probe measurements reduce MAE by 86%.

Connection to physics. The Gumbel race is the same mechanism underlying EVT in physics. Synthetic validation with equicorrelated Gaussian data ($K \in \{4, 10, 20\}$, $d = 50$, generator parameter $\rho \in [0, 0.8]$) supports the logit-linear form (mean $R^2 = 0.982$). The fitted α showed no systematic dependence on the generator’s ρ , so this experiment provides no support for an $\alpha(\rho)$ prediction. The two-regime alpha structure ($\alpha \approx 1.5$ for discrete tokens vs. $\alpha \approx 3.3\text{--}5.0$ for continuous signals) remains an empirical observation; its mechanism is open, and the retracted equicorrelation analysis cannot explain it.

7 Limitations

1. **Cross-task drift.** LODO CVs (0.419, 0.920) confirm constants are task-specific. One calibration measurement per dataset is required; ranking within a dataset is unaffected.
2. **External replication needed.** Results are internal; independent reproduction is needed.
3. **Large-K attenuation.** Both ViT and CNN at $K = 100$ show $r \approx 0.75$, below the MC noise floor ($E[r] = 0.875$). Hierarchical class structure creates non-ideal kappa geometry.
4. **Assumption dependence.** The proof relies on concentration and shared-scale Gumbel approximations; tighter bounds remain open.
5. **Protein language models.** Seven protein LMs from 3 architecture families (ESM-2 at 4 sizes, ProtBERT, ProtBERT-BFD, ESM-1b; EC-number classification, $K = 7$) yield $\alpha = -1.17$, $r = -0.15$, $p = 0.76$: the law *fails* across protein architectures (supplementary materials). EC classes are extremely heterogeneous, so protein LMs achieve accuracy via within-class compactness rather than global centroid separation. The law requires semantically coherent categories.
6. **Human fMRI.** Eleven visual ROIs from the Natural Scenes Dataset (NSD, subject 01, $K = 12$ COCO supercategories) yield chance-level 1-NN accuracy; pooled across all ROIs and classes ($n = 132$), $r = 0.12$, $p = 0.18$, non-significant (supplementary materials). Raw BOLD voxel responses lack sufficient SNR for nearest-neighbor classification. The law is validated on high-SNR neural measurements (spike trains, multi-electrode arrays) but not raw BOLD.

Extended limitations (biological scope, multilabel regime, surgery scope, beyond-1NN, two-component geometry) are detailed in Appendix J.

8 Conclusion

We present the CTI Universal Law, derived from extreme-value theory, with functional form proven as a conditional theorem. The law is validated across 12 NLP architectures (LOAO $\text{CV}_\alpha = 0.023$), in blind OOD tests ($r = 0.817$), an expanded holdout (11 models $\times$ 8 datasets, all 6 criteria pass), cross-modally (ViT $R^2 = 0.964$, ResNet50 $r = 0.749$, 7 speech models $r = 0.898$ with $\alpha_{\text{audio}} = 4.67$), and in biological visual cortex (30/32 sessions). The Gumbel-race mechanism extends to generation ($r = -0.55$, 22 models; $r = -0.70$ excluding one outlier). Universality is of the *functional form*: the constant α is signal-type-specific

($\alpha_{\text{NLP}} \approx 1.48$ for discrete tokens; $\alpha_{\text{ViT}} \approx 4.5$, $\alpha_{\text{CNN}} \approx 4.0$, $\alpha_{\text{audio}} \approx 4.7$ for continuous signals), with all continuous-signal modalities converging at $\alpha \approx 3.3\text{--}5.0$ and discrete-token NLP at $\alpha \approx 1.5$; honest-scope limitations including protein LM failure and human fMRI null are enumerated in §7. The formerly reported cross-modal equicorrelation constant, near-simplex universality, and $\alpha(\rho)$ explanation are retracted rather than retained as approximate claims. Open directions include a renormalization theory for cross-family constants and independent replication. Independent convergent evidence: St-Yves et al. (2025) confirm centroid-separation geometry predicts decoding in human visual cortex.

Reproducibility Statement

All numerical claims are taken directly from repository JSON artifacts. All pre-registrations were committed before data collection; timestamps are verifiable upon code release. Code and data will be released upon acceptance.

References

- J. Kaplan, S. McCandlish, T. Henighan, et al. Scaling laws for neural language models. *arXiv:2001.08361*, 2020.
- J. Hoffmann, S. Borgeaud, A. Mensch, et al. Training compute-optimal large language models. *NeurIPS*, 35:30016–30030, 2022.
- R. A. Fisher and L. H. C. Tippett. Limiting forms of the frequency distribution of the largest or smallest member of a sample. *Proceedings of the Cambridge Philosophical Society*, 24:180–190, 1928.
- B. Gnedenko. Sur la distribution limite du terme maximum d’une serie aleatoire. *Annals of Mathematics*, 44(3):423–453, 1943.
- E. J. Gumbel. *Statistics of Extremes*. Columbia University Press, 1958.
- Z. Wu, Y. Xiong, S. X. Yu, and D. Lin. Unsupervised feature learning via non-parametric instance discrimination. *CVPR*, 2018.
- M. Caron, H. Touvron, I. Misra, et al. Emerging properties in self-supervised vision transformers. *ICCV*, 2021.
- V. Pappayan, X. Y. Han, and D. L. Donoho. Prevalence of neural collapse during the terminal phase of deep learning training. *PNAS*, 117(40):24652–24663, 2020.
- X. Y. Han, V. Pappayan, and D. L. Donoho. Neural collapse under cross-entropy losses. *PNAS*, 119(6):e2111457119, 2022.
- S. Biderman, H. Schoelkopf, Q. G. Anthony, et al. Pythia: A suite for analyzing large language models across training and scaling. *ICML*, 2023.
- X. Zhang, J. Zhao, and Y. LeCun. Character-level convolutional networks for text classification. *NeurIPS*, 2015.
- J. Lehmann, R. Isele, M. Jakob, et al. DBpedia – a large-scale, multilingual knowledge base extracted from Wikipedia. *Semantic Web*, 6(2):167–195, 2015.
- K. Lang. Newsweeder: Learning to filter netnews. *ICML*, 1995.
- D. Demszky, D. Movshovitz-Attias, J. Ko, et al. GoEmotions: A dataset of fine-grained emotions. *ACL*, 2020.
- A. Dosovitskiy, L. Beyer, A. Kolesnikov, et al. An image is worth 16x16 words: Transformers for image recognition at scale. *ICLR*, 2021.
- Q. Garrido, A. Balestriero, L. Najman, and Y. LeCun. RankMe: Assessing the downstream performance of pretrained self-supervised representations by their rank. *ICML*, 2023.

276 Z. Zhu, T. Ding, J. Zhou, X. Li, C. You, J. Sulam, and Q. Qu. A geometric analysis of neural
277 collapse with unconstrained features. *NeurIPS*, 2021.

278 A. Rangamani, M. Lindegaard, T. Galanti, and T. Poggio. Feature learning in deep classifiers
279 through intermediate neural collapse. *ICML*, 2023.

280 J. Liu, A. Koppara, and S. Levine. Tracing the representation geometry of language models
281 from pretraining to post-training. *arXiv:2509.23024*, 2025.

282 J. H. Siegle, X. Jia, S. Durand, S. Gale, C. Bennett, N. Graddis, G. Bhaskaran, and C. Koch.
283 Survey of spiking in the mouse visual system reveals functional hierarchy. *Nature*,
284 592(7852):86–92, 2021.

285 G. St-Yves, K. Kay, and T. Naselaris. Variation in the geometry of concept manifolds across
286 human visual cortex. *PLOS Computational Biology*, 21(9):e1013416, 2025.

287 D. A. Pospisil and J. W. Pillow. Revisiting the high-dimensional geometry of population
288 responses in the visual cortex. *PNAS*, 122, 2025.

289 Z. Chen and M. F. Bonner. Universal dimensions of visual representation. *Science Advances*,
290 11(27):eadw7697, 2025.

291 M. Munn, B. Dherin, and J. Gonzalvo. The impact of geometric complexity on neural
292 collapse in transfer learning. *NeurIPS*, 2024.

293 R. Wu and V. Papyan. Linguistic collapse: Neural collapse in (large) language models.
294 *NeurIPS*, 2024.

295 A. Kulkarni, J. Springer, A. Subramonian, and S. Swayamdipta. Disentangling geometry,
296 performance, and training in language models. *arXiv:2602.20433*, 2026.

297 Z. Yang, Z. Dai, R. Salakhutdinov, and W. W. Cohen. Breaking the softmax bottleneck: A
298 high-rank RNN language model. *ICLR*, 2018.

299 J. Ye, J. Hong, and M. Song. Multiplicative logit adjustment approximates neural-collapse-
300 aware decision boundary adjustment. *arXiv:2409.17582*, 2024.

301 M. Y. Harun, J. Gallardo, and C. Kanan. Controlling neural collapse enhances out-of-
302 distribution detection and transfer learning. *ICML*, 2025.

303 J. Zhao, T. S. Cheng, W. Masarczyk, and A. Lucchi. Optimizer choice matters for the
304 emergence of neural collapse. *ICLR*, 2026.

305 C. F. Cadieu, H. Hong, D. L. K. Yamins, N. Pinto, D. Ardila, E. A. Solomon, N. J. Majaj, and
306 J. J. DiCarlo. Deep neural networks rival the representation of primate IT cortex for core
307 visual object recognition. *PLOS Computational Biology*, 10(12):e1003963, 2014.

308 C. Stringer, M. Pachitariu, N. Steinmetz, C. B. Reddy, M. Carandini, and K. D. Harris.
309 Spontaneous behaviors drive multidimensional, brainwide activity. *Science*, 364(6437):255,
310 2019.

311 Z. Zhang, F. Lin, H. Liu, J. Morales, H. Zhang, K. Yamada, V. B. Kolachalama, and
312 V. Saligrama. GPS: A probabilistic distributional similarity with Gumbel priors for
313 set-to-set matching. *ICLR*, 2025.

314 N. Lourie, M. Y. Hu, and K. Cho. Scaling laws are unreliable for downstream tasks: A reality
315 check. *Findings of EMNLP*, 2025.

316 R. Jha, C. Zhang, V. Shmatikov, and J. X. Morris. Harnessing the universal geometry of
317 embeddings. *arXiv:2505.12540*, 2025.

318 P. Súkeník, C. H. Lampert, and M. Mondelli. Neural collapse is globally optimal in deep
319 regularized ResNets and Transformers. *arXiv:2505.15239*, 2025.

A Derivation Details

Starting from the Gumbel race identity (Eq. 3), under nearest-class dominance: $q \approx [1 + (K - 1) \exp(-\Delta_k/s)]^{-1}$. Taking logits: $\text{logit}(q) = \Delta_k/s - \log(K - 1)$. With $\Delta_k = a_0 + a_1 \kappa_{\text{nearest}} + o(1)$: $\text{logit}(q) = \underbrace{a_1/s \kappa_{\text{nearest}}}_{\alpha} - \underbrace{1}_{\beta_{\infty}} \log(K - 1) + \underbrace{a_0/s}_{C_0} + o(1)$. Finite-data corrections renormalize $\beta_{\infty} = 1$ to $\beta \approx 0.746$.

Retracted equicorrelation analysis. For provenance, the legacy estimator reported six cross-modal values in $[0.455, 0.467]$ with mean 0.462 (CV= 1.0%; results/cti_cross_modal_rho.json). The code multiplied projected centroid differences by $\sqrt{\Lambda}$, applying $\Sigma_W^{1/2}$ rather than the stated inverse-whitening transform $\Sigma_W^{-1/2}$. Matched labeled-cloud nulls place the historical band inside central 99% null intervals, and a shared-anchor isotropic-centroid calculation yields $\rho \rightarrow 1/2$ without learned or semantic structure. No corrected real-data comparison is available because the retained caches lack aligned embeddings and labels. Accordingly, the historical values are retained only as measurements from a faulty estimator; they do not support universal or approximately universal equicorrelation, near-simplex geometry, or a zero-parameter $\alpha(\rho)$ explanation.

Hierarchical beta derivation. Consider K classes organized into $\sqrt{K}$ groups of $\sqrt{K}$ each, with intra-group distance D_1 and inter-group distance $D_2 \gg D_1$. The Gumbel race sum decomposes: when $D_2 - D_1 \gg s$, the across-group term is negligible and $N_{\text{eff}} \approx \sqrt{K} - 1$, giving $\beta = 0.5$. This matches the empirical $\hat{\beta} = 0.478$.

B Pre-registration Deviation Log

Three post-registration procedural edits were required: (1) dataset loading API deprecation (TREC→emotion, identical K); (2) numeric stability patch for rare classes; (3) RWKV interval documentation error. None changes predictions, thresholds, or metrics. Full table with commit hashes available in the extended version.

C Surgery Details

Pre-registered global surgery (CIFAR-100, $K = 20$, linear regime, 3 seeds): $\Delta \text{logit}_{\text{global}} / \Delta \text{logit}_{\text{single}} = 18.74 \approx d_{\text{eff}} = 24.52$. Cross-domain NLP follow-ups failed quantitatively (ratio $\approx 2\text{--}4 \ll d_{\text{eff}}$); directional ordering confirmed. The $1/d_{\text{eff}}$ mechanism appears CNN-specific.

D Competition Field: Extended Results

The full competition field analysis includes: (1) 6-arm RCT confirming j_2 is causally active ($r = 0.811$); (2) soft-min competition field $\phi(\tau^*)$ with $\tau^* = 0.20$ improving held-out R^2 from 0.451 to 0.578 (+28%); (3) additivity RCT ($r = 0.9964$); (4) confusion-matrix Gumbel test (pooled $r = 0.750$, $\rho = 0.768$); (5) unconfounded weight map ($\sum_r w_r = 2.42$, heavy-head decay); (6) cross-architecture competition scaling (5/5 pass r-scale ≥ 0.90).

Causal confusion-matrix prediction. With fixed $A = 1.054$ and $\tau^* = 0.20$ (no refit), the law predicts how every confusion-matrix entry changes under centroid shifts. At $\delta = 1$: $r = 0.842$, sign acc. 92.9%. At $\delta = 2$: $r = 0.808$, sign acc. 100%. At $\delta = 3$: $r = 0.776$, sign acc. 100%. All $p < 10^{-35}$, $n = 182$ class pairs per level.

Top- m competitor sweep. The fitted slope rises from $\alpha_1 = 1.052$ ($m = 1$) to $\alpha_2 = 1.523$ ($m = 2$; 2.3% from LOAO $\alpha = 1.477$) to $\alpha_{13} = 3.015$ ($m = K - 1$), confirming exactly 2 effective competitors drive the cross-architecture constant.

E Sparse Competition

Pre-registered constrained comparison on 444 points (19 architectures, 10 datasets): M2 ($\beta = 0.5$) beats M1 ($\beta = 1.0$) by 10.1% LOAO MAE and 24.8% LODO MAE. Unconstrained MLE: $\hat{\beta} = 0.478$ (within 0.6% of $\beta = 0.5$). Triangulated by synthetic sweep ($B \propto \sqrt{K-1}$, CV=4.4%) and direct centroid-geometry N_{eff} scaling (mean exponent 0.84, $n = 7$ architectures).

F Comprehensive Universality

A 19-architecture $\times$ 10-dataset stress test ($K = 4-77$, 444 points): LOAO $\text{CV}_\alpha = 1.75\%$ (PASS $< 5\%$), LODO $\text{CV}_\alpha = 5.51\%$ (PASS $< 7\%$), median per- K $r = 0.770$. Kappa spread (17 $\times$ more predictive than K ; $p = 0.038$) is the mechanistic predictor of ranking reliability.

K-scaling. $A(K) = 1.164 / \log K + 1.301$ ($r = 0.93$). One-point calibration: $\text{MAE} \leq 0.08$ on all 6 prospective datasets. H3 ranking: 9 architectures on Banking77, Spearman $\rho = 0.833$, $p = 0.005$ (PASS).

G H8+ Expanded Holdout Detail

Table 4: H8+ per-model holdout (11 unseen architectures, all parameters frozen).

Model	Family	n	MAE	Max err	Status
distilbert-base	encoder	8	0.036	0.060	Best
albert-base-v2	encoder	7	0.067	0.122	Good
roberta-base	encoder	8	0.071	0.309	Good
falcon-rw-1b	decoder	8	0.051	0.110	Good
stablelm-3b-4e1t	decoder	8	0.070	0.155	Good
opt-125m	decoder	7	0.071	0.148	Good
qwen2.5-1.5b	decoder	7	0.077	0.194	Good
phi-1.5	decoder	8	0.082	0.183	Good
pythia-2.8b	decoder	8	0.082	0.172	Good
bloom-560m	decoder	4	0.160	0.316	Weak
gemma-3-1b	decoder	4	0.167	0.533	Weak
All (full H8+)		77	0.077	0.533	PASS

H Two-Component Geometry

Null-space identifiability: q tracks κ_{nearest} (OLMo-1B: $r = 1.00$), not κ_{eff} . 2D causal surface: bivariate model does not outperform univariate ($\Delta R^2 \leq 0.03$). Two-knob identifiability (6/7 criteria fail): κ_{nearest} alone is the causal driver.

I Generation Law: Model Details

J Extended Limitations

- Causal controls.** Some farthest-arm effects are non-zero; intervention geometry can leak.
- Surgery scope.** Global surgery matches prediction in CIFAR-100 linear regime; NLP follow-ups fail quantitatively.
- Ceiling effects.** High-accuracy regimes flatten logits and attenuate slopes.
- Small-K regime** ($K \leq 3$). Functional form holds ($r = 0.42-0.52$) but magnitude deviates $\sim 70\%$.
- Multilabel.** go_emotions ($K = 28$, multilabel) fails; the law applies to single-label only.

Table 5: Generation law: 22-model suite ordered by CE.

Model	Arch	N (M)	V	κ_{top1K}	CE
Mistral-7B	Trans.	7000	32K	0.876	1.551
Falcon-H1-3B	Hybrid	3000	131K	0.949	1.564
Falcon-H1-1.5B	Hybrid	1500	131K	0.977	1.741
Falcon-H1-0.5B	Hybrid	500	131K	0.716	1.935
Mamba-2.8B	SSM	2800	50K	0.720	2.110
Pythia-2.8B	Trans.	2800	50K	0.794	2.174
SmolLM2-360M	Trans.	360	49K	0.725	2.179
Granite-Micro	Hybrid	500	49K	0.773	2.183
Mamba-1.4B	SSM	1400	50K	0.720	2.218
Qwen2-0.5B	Trans.	500	152K	0.716	2.292
Pythia-1.4B	Trans.	1400	50K	0.795	2.292
Mamba-790M	SSM	790	50K	0.724	2.308
Pythia-1B	Trans.	1000	50K	0.821	2.398
Qwen3-4B	Trans.	4000	152K	0.710	2.422
Mamba-370M	SSM	370	50K	0.711	2.455
Pythia-410M	Trans.	410	50K	0.791	2.565
Qwen3-1.7B	Trans.	1700	152K	0.698	2.588
Qwen3-0.6B	Trans.	600	152K	0.673	2.777
Mamba-130M	SSM	130	50K	0.697	2.808
Pythia-160M	Trans.	160	50K	0.251	3.093
GPT-2	Trans.	124	50K	0.730	3.108
LFM2.5-1.2B	Novel	1200	66K	0.863	3.349

10. **Biological constants.** Law form confirmed; constants are $15\text{--}34\times$ smaller than NLP. Stronger NC can hurt generalization (Harun et al., 2025); κ_{nearest} maximization may not be globally optimal.
11. **Beyond-1NN.** κ_{nearest} predicts $k = 3, 5$ NN ($r \approx 0.79$); linear probe prediction fails within-dataset ($r = 0.22\text{--}0.59$). H3 ranking at $n = 9$: $\rho = 0.833$, $p = 0.005$ (PASS).
12. **Optimizer documentation.** All tested architectures use Adam-family optimizers; Zhao et al. (2026) shows optimizer affects NC emergence.
13. **Protein language models.** Seven protein LMs from 3 families (ESM-2, ProtBERT, ESM-1b; EC classification $K = 7$): $\alpha = -1.17$, $r = -0.15$, $p = 0.76$ — law fails. EC classes span thousands of structurally diverse enzymes; centroid separation is not the classification mechanism.
14. **Human fMRI.** NSD (11 visual ROIs, $K = 12$ COCO supercategories): chance-level 1-NN ($q \approx 0.08$); pooled $r = 0.12$, $p = 0.18$. Raw BOLD lacks SNR for nearest-neighbor classification; the law requires measurable above-chance 1-NN accuracy.

K Biological Details

Macaque IT (Cadieu et al., 2014) ($K = 7$): per-image $r = 0.41$; V4 $r = 0.116$ (hierarchy gradient). Mouse V1 (Stringer et al., 2019) ($K = 11$): per-image $r = 0.64$. Allen Neuropixels: $K = 118$, 32 mice, response window 50–250ms. Population geometry dominated by ≈ 10 effective dimensions (Pospisil and Pillow, 2025). Universal dimensions: Chen and Bonner (2025) show diverse networks converge on < 10 latent dimensions aligned with brain, consistent with $d_{\text{eff,comp}} \approx 2$.

L Scaling Dynamics

Power-law fit $\kappa \propto N^\gamma$: $\hat{\gamma} = 0.003$, $R^2 = 0.008$, $p = 0.91$ — κ_{nearest} does not scale with model size. MAP@10 does scale ($R^2 = 0.90$, $p = 0.05$).